

SPRING 2019

CS464 Final Report

Introduction to Machine Learning

Selin Özdaş 21400537

Büşra Arabacı 21401688

Cem Demirpek 21301643

Doğan Can Eren 21400329

Elif Engin



Introduction

The drastic increase in the number of artworks that are digitized created a data pool. This pool required people to find ways to retrieve and archive the data. Artists use different concepts to describe paintings such as texture, colors, patterns and balance. The main problem caused by this situation is that some features may give high variations. Mona is a categorization tool of fine art paintings in terms of their specifications such as genre, style etc. by processing visual features. To demonstrate Mona, we are going to use a publicly available dataset from Kaggle which is [Painter by Numbers](#) prepared by Kiri Nichol.

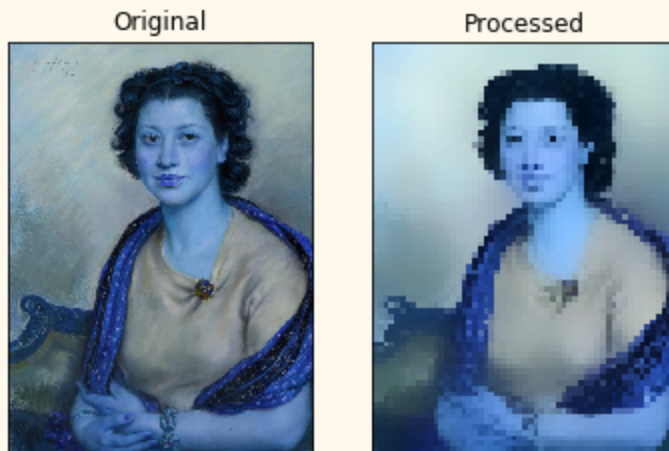
Background Information

"According to which metrics does an art historian categorizes the artwork?" is one of the major questions arose after the institutions of art gain power[2]. While there were subjective opinions about how art historians do that, emergence of new technologies enhanced our understanding of artworks. Art historians turned image processing techniques to good account to get a better understanding of artworks. It is found out that we can encode information of fine art paintings by applying unsupervised and supervised learning methodologies[3]. Additionally, discovery of convolutional neural network(CNN) modelling was a breakthrough in image classifications tasks. It turned out that CNN modelling is beneficial due to the decrease in parameters and the reduction in the amount of training data to save time. It was like seeing the greater picture by using a smaller frame.

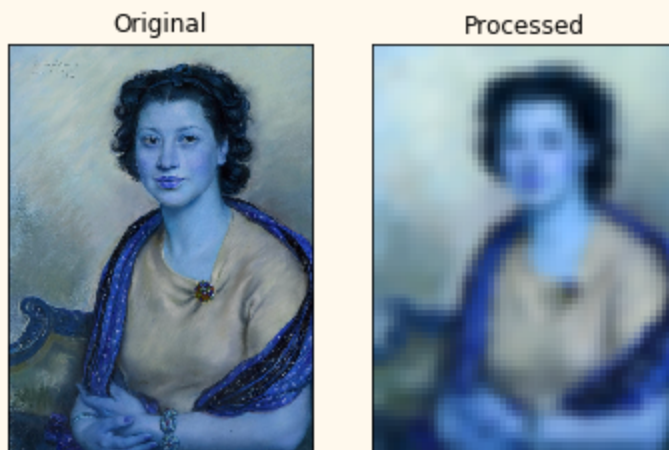
Work Done

Preprocessing images

In order to feed the images to machine learning and deep learning models, some preprocessing needed to be done because images in the dataset varies in size and color balance. Image sizes were changed and all images were converted to same size. Also, rgb color values are normalized to the range between 0 and 1 because using other values causes neural networks to converge to 1 for all outputs since the input values are high which affect the classification. Moreover, gaussian blur filter was applied to images to remove noise. Bilateral image filtering was tried and the errors could not resolved, however it looks promising as it produces a clean image for training:



Compared to Gaussian filtering:



Opencv library is used for image resizing and noise reduction. Additionally, some data cleaning was done in the dataset because dataset includes some corrupted images and images with mode different than rgb.

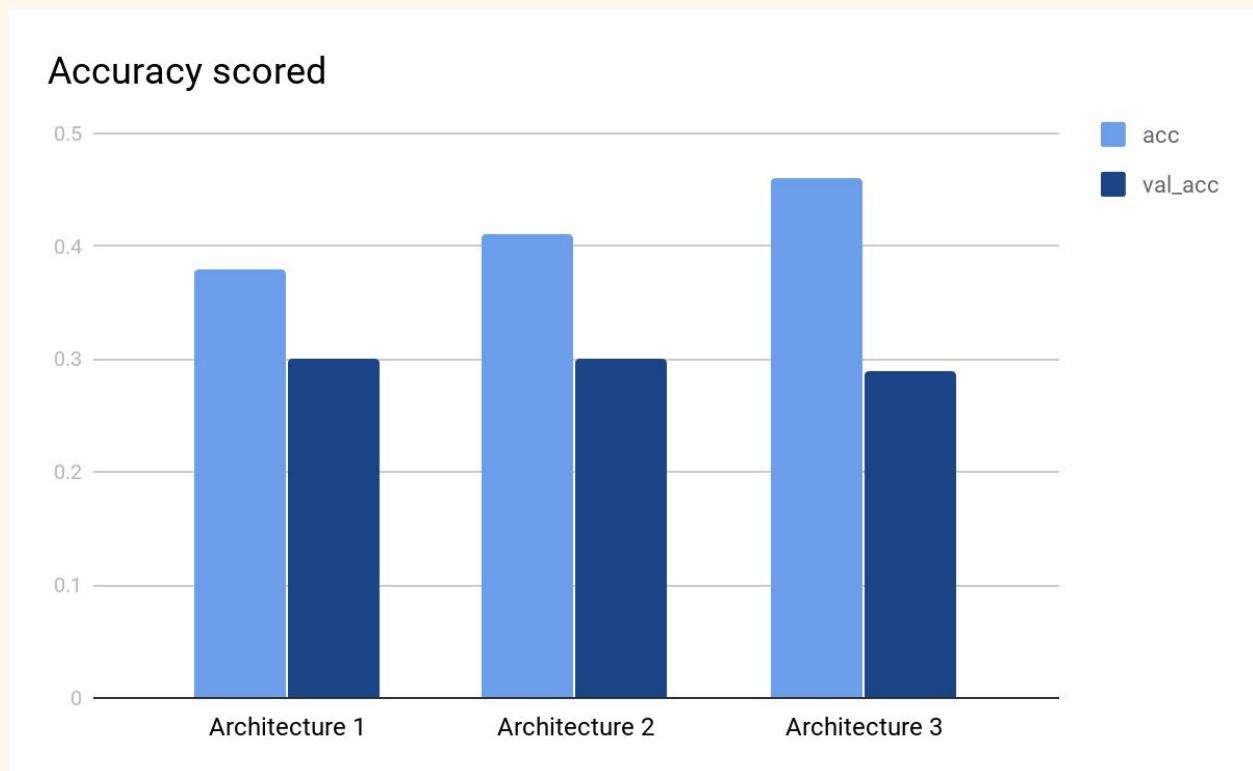
Image classification

Our focus for this iteration is to classify images according to their genre and style. Artificial neural networks and convolutional neural networks were trained for this purpose. Since they are both multiclass classification, we only focused on genre and in the next iteration, we will apply the best model for style classification.

We had 8241 samples in the dataset after data cleaning. For the genre and style classification, since we have 127 styles and 40 genres and the task becomes

multiclass classification. Since it is multiclass classification, we encoded the classes by using one hot encoding technique. Therefore, for genre classification, we had one hot vectors of length 40 and for style we had one hot vectors of length 127 for labels.

We split the dataset into train and test set. We used $\frac{1}{3}$ of the data as test set. We first trained artificial neural network models for creating baseline model. We tried different architectures by changing hidden layer number and number of neurons in the layers. Also, we experimented with different optimizers. For these neural networks, stochastic gradient descent algorithm is used with batch size 1. Results with artificial neural networks:



Architecture: 100*100*3 -> 256 -> 64 -> 40 Results: acc:0.38, val_acc: 0.30

Architecture: 100*100*3 -> 1024 -> 256 -> 64 -> 40 Results: acc: 0.41 , val_acc: 0.30

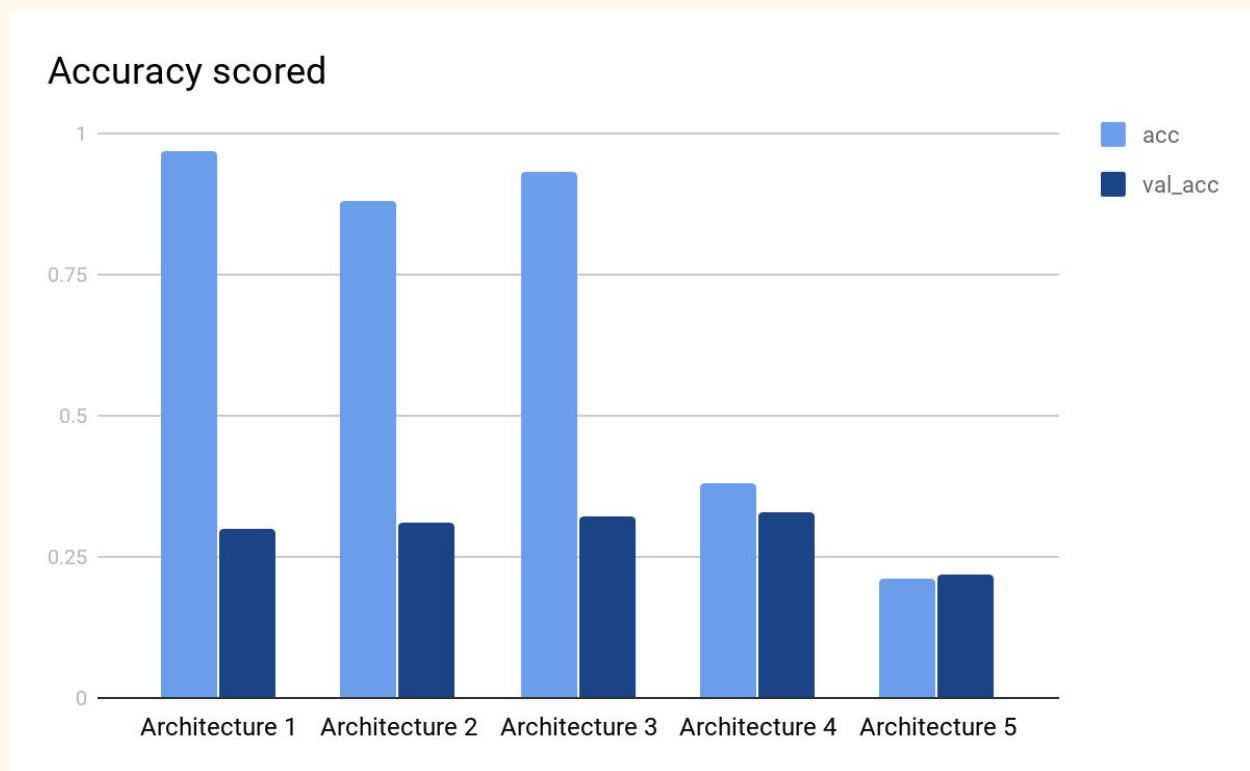
Architecture: 100*100*3 -> 3000 -> 40 Results: acc: 0.46 , val_acc: 0.29

According to artificial neural network results, network with one hidden layer and large hidden layer size performs best on training set. But their results on test set are similar. As we can see these architectures are not powerful enough to capture images features. They give low accuracies because they are simple neural

network architectures. Therefore, we tried Convolutional Neural Network architectures for image classification task.

CNN Models

With Convolutional Neural Networks, we experimented with different number of convolutional layers, different number of hidden layers. Also, we experimented with adding max pooling option but we realised that it does not help in our case. Pooling helps on increasing the receptive field of a filter, which is defined as the region in the input space that a particular Convolutional Neural Network's feature is looking at[1]. It would help if we were looking for a specific region in the input like eyes, but as we are not; it reduces our model's accuracy.



Architecture: kernelsize : 9, conv32 -> conv32-> conv64-> conv128 -> 512 -> 128 -> 40

Results: acc:0.97, val_acc: 0.30

Architecture: kernelsize : 3, conv32 -> conv64-> conv64-> conv96 -> conv32->128->40

Results: acc:0.88, val_acc: 0.31

Architecture: kernelsize : 3, conv32 -> conv64-> conv64-> conv32 ->128->40

Results: acc:0.93, val_acc: 0.32

Architecture: kernelsize : 3, conv32 -> conv64-> conv64-> conv128-> 256->128->40

Results: acc: 0.38, val_acc: 0.33

Architecture: kernelsize : 5, conv32 -> conv64-> conv64-> conv96 -> conv32->128->40

Results: acc:0.21, va_acc: 0.22

Architecture: kernelsize : 5, conv96 -> conv128-> conv128-> conv96 -> conv64->512->40 Results: acc:0.92, va_acc: 0.29

For the convolutional neural networks, we obtained good results for the training set. However, it gives bad accuracy for test set. This is a sign for overfitting. There are several ways to solve overfitting problem and one of them is adding regularization to the model. We experimented on the best result we get by increasing dropout rates by .25, and we get the following results:

Architecture: kernelsize : 3, conv32 -> conv64-> conv128-> 256->128->40

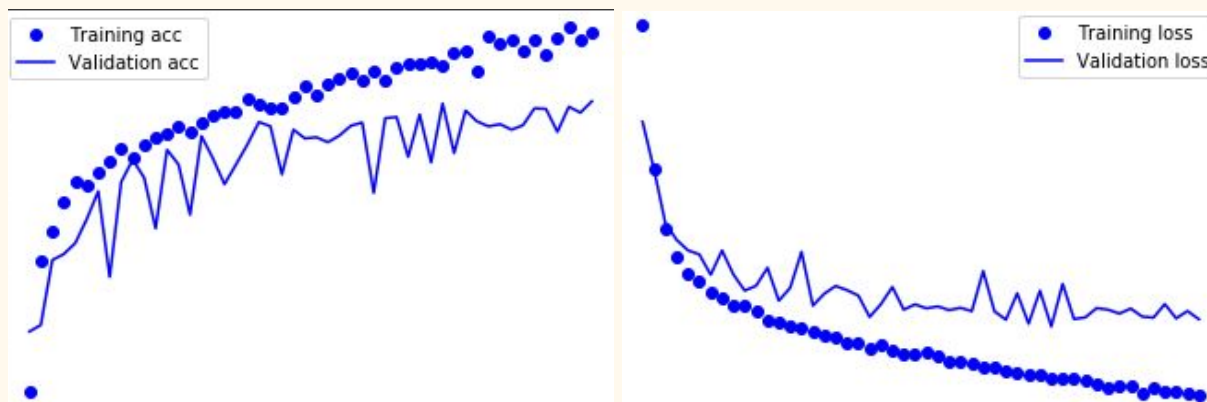
Results: acc: 0.33, va_acc: 0.35

Final CNN Model

Model 1

This model includes 3 convolutional layer followed by 2 fully-connected layer. After each layer (except the last) there is 0.5 dropout rate. Between the convolutional layers there are three max-pooling layers with 2 x 2 window size. The preferred optimizer is Adam.

Results = Epoch 50 loss: 1.9588 - acc: 0.4075 - va_loss: 2.2906 - va_acc: 0.3610



Model Summary =

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 48, 48, 32)	896
batch_normalization_7 (Batch Normalization)	(None, 48, 48, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 24, 24, 32)	0

dropout_7 (Dropout)	(None, 24, 24, 32)	0
conv2d_6 (Conv2D)	(None, 22, 22, 32)	9248
max_pooling2d_2 (MaxPooling2)	(None, 11, 11, 32)	0
dropout_8 (Dropout)	(None, 11, 11, 32)	0
conv2d_7 (Conv2D)	(None, 9, 9, 64)	18496
batch_normalization_8 (Batch Normalization)	(None, 9, 9, 64)	256
max_pooling2d_3 (MaxPooling2)	(None, 4, 4, 64)	0
dropout_9 (Dropout)	(None, 4, 4, 64)	0
flatten_2 (Flatten)	(None, 1024)	0
dense_4 (Dense)	(None, 64)	65600
dropout_10 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 40)	2600

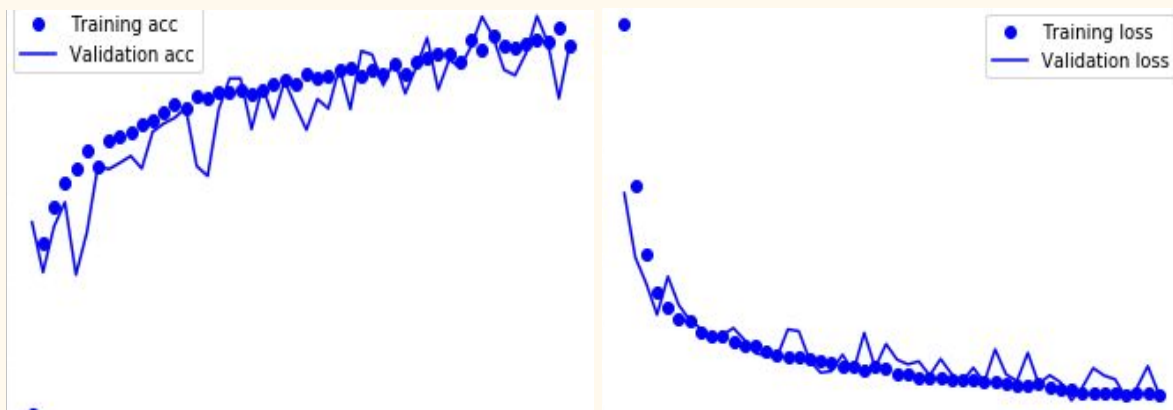
=====

Total params: 97,224
Trainable params: 97,032
Non-trainable params: 192

Model 2

As we overfit to the data, to decrease overfitting we changed the max pooling window size to be 3 x 3 for the first two max-pooling layers and decreased the overfitting problem and the loss without considerable amount of decrease in accuracy.

Results = Epoch 47 loss: 2.2842 - acc: 0.3396 - val_loss: 2.2750 - val_acc: 0.3540



Model Summary =

Layer (type)	Output Shape	Param #
=====		
conv2d_23 (Conv2D)	(None, 48, 48, 32)	896
max_pooling2d_16 (MaxPooling)	(None, 16, 16, 32)	0
batch_normalization_23 (Batch Normalization)	(None, 16, 16, 32)	128
dropout_29 (Dropout)	(None, 16, 16, 32)	0
conv2d_24 (Conv2D)	(None, 14, 14, 32)	9248
max_pooling2d_17 (MaxPooling)	(None, 4, 4, 32)	0
batch_normalization_24 (Batch Normalization)	(None, 4, 4, 32)	128
dropout_30 (Dropout)	(None, 4, 4, 32)	0
conv2d_25 (Conv2D)	(None, 2, 2, 64)	18496
max_pooling2d_18 (MaxPooling)	(None, 1, 1, 64)	0
batch_normalization_25 (Batch Normalization)	(None, 1, 1, 64)	256

dropout_31 (Dropout)	(None, 1, 1, 64)	0
flatten_7 (Flatten)	(None, 64)	0
dense_14 (Dense)	(None, 64)	4160
batch_normalization_26 (Batch Normalization)	(None, 64)	256
dropout_32 (Dropout)	(None, 64)	0
dense_15 (Dense)	(None, 40)	2600

=====

Total params: 36,168
Trainable params: 35,784
Non-trainable params: 384

Model 3

We added L1 regularization to the data to make it better. Even though it did not help us to increase our accuracy, now our validation accuracy becomes more than the training accuracy

Results = Epoch 48 loss: 2.2435 - acc: 0.3299 - val_loss: 2.2247 - val_acc: 0.3505

Model Summary:

Layer (type)	Output Shape	Param #
conv2d_150 (Conv2D)	(None, 48, 48, 32)	896
max_pooling2d_67 (MaxPooling)	(None, 16, 16, 32)	0
batch_normalization_150 (Batch Normalization)	(None, 16, 16, 32)	128
dropout_152 (Dropout)	(None, 16, 16, 32)	0
conv2d_151 (Conv2D)	(None, 14, 14, 32)	9248
max_pooling2d_68 (MaxPooling)	(None, 4, 4, 32)	0

batch_normalization_151 (Batch Normalization)	(None, 4, 4, 32)	128
dropout_153 (Dropout)	(None, 4, 4, 32)	0
conv2d_152 (Conv2D)	(None, 2, 2, 64)	18496
max_pooling2d_69 (MaxPooling2D)	(None, 1, 1, 64)	0
batch_normalization_152 (Batch Normalization)	(None, 1, 1, 64)	256
dropout_154 (Dropout)	(None, 1, 1, 64)	0
flatten_28 (Flatten)	(None, 64)	0
dense_55 (Dense)	(None, 32)	2080
batch_normalization_153 (Batch Normalization)	(None, 32)	128
dropout_155 (Dropout)	(None, 32)	0
dense_56 (Dense)	(None, 40)	1320

=====
 Total params: 32,680
 Trainable params: 32,360
 Non-trainable params: 320



Results & Conclusion

Since we are doing multiclass classification and our class size is large, we could not get high accuracies. Although convolutional neural networks generally perform good in image classification tasks, they are not enough for large classes. Also, our dataset size is not enough. We think that our models cannot see enough examples for classification from this dataset. Our models would work better with more data and better regularization techniques. Most articles say that going deeper than going wider gives best results. Even though we have experimented with too many models, we could not see this effect on our accuracy.

References

- [1] D. H. T. Hien and D. H. T. Hien, "A guide to receptive field arithmetic for Convolutional Neural Networks," *Medium*, 05-Apr-2017. [Online]. Available: <https://medium.com/mlreview/a-guide-to-receptive-field-arithmetic-for-convolutional-neural-networks-e0f514068807>. [Accessed: 05-Apr-2019].
- [2] Stam, Deirdre Corcoran. "HOW ART HISTORIANS LOOK FOR INFORMATION." *Art Documentation: Journal of the Art Libraries Society of North America*, vol. 3, no. 4, 1984, pp. 117-119. JSTOR, www.jstor.org/stable/27947353.
- [3] T. E. Lombardi. "The classification of style in fine-art painting". ETD Collection for Pace University. Paper AAI3189084., 2005.
- [4] K. Sorokina, "Image Classification with Convolutional Neural Networks", *Medium*, 2019. [Online]. Available: <https://medium.com/@ksusorokina/image-classification-with-convolutional-neural-networks-496815db12a8>.